

Database Management System (DBMS)

—Vishvadeep Yothi playlist

Database :- The collection of data, usually referred to as the database, contains information relevant to an enterprise.

Database Management System :- The database management system (DBMS) is a collection of interrelated data and a set of programs to access those data.

↳ info retrieval,
insertion, deletion,
update.

Goal of DBMS

Providing a way to store & retrieve database info
that is both convenient & efficient.

Ensuring safety of information

Disadvantages of file system

- Data Redundancy and Inconsistency
- Difficulty in Accessing Data
- Data isolation
- Integrity problems (applying constraints)
- Atomicity problem
- Concurrent - Access Anomalies
- Security problems

Instance - The collection of information stored in the database at a particular moment is called an instance of database.

Schema - The overall design of the database is called database schema.

Database Languages

Data Definition Language (DDL)

(schema related changes are done using DDL)

Data Manipulation Language (DML)

(data related changes are done using DML)

Procedural DMLs

requires user to specify what data is needed & how to get that data

Non-procedural (Declarative DMLs)

requires user to specify which data is needed without specifying how to get those data

View of data -

1. Physical level -

how data is actually stored in physical memory.

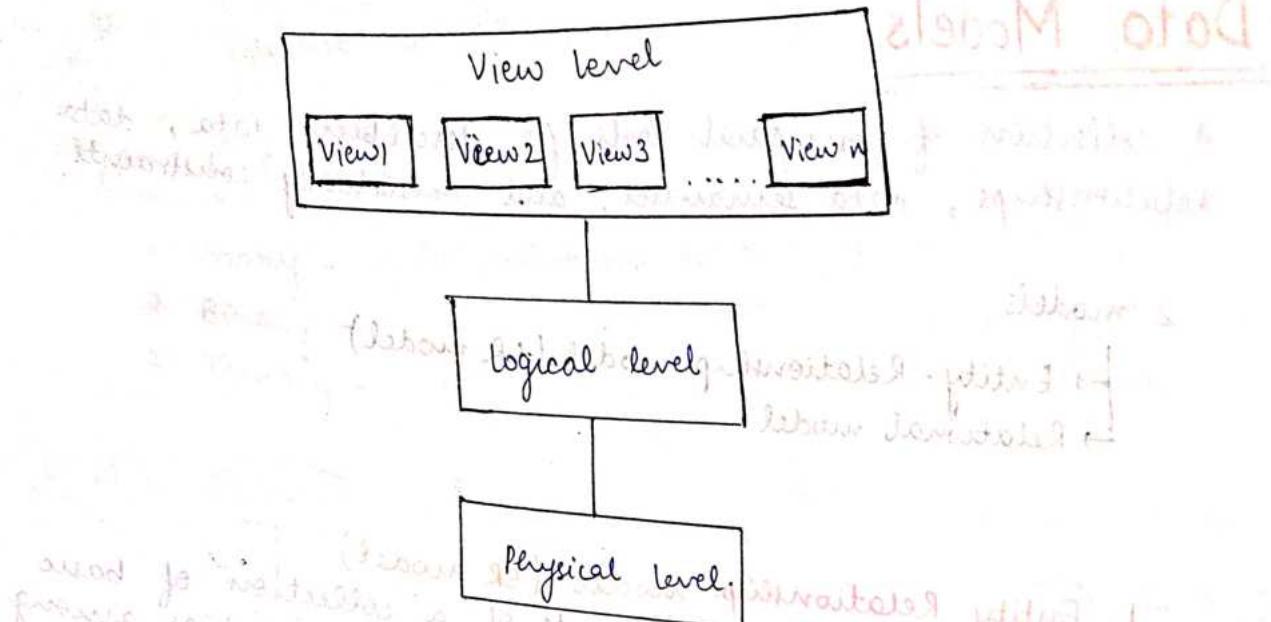
2. Logical level -

tables, at this level, we see data distributed in the form of tables in database.

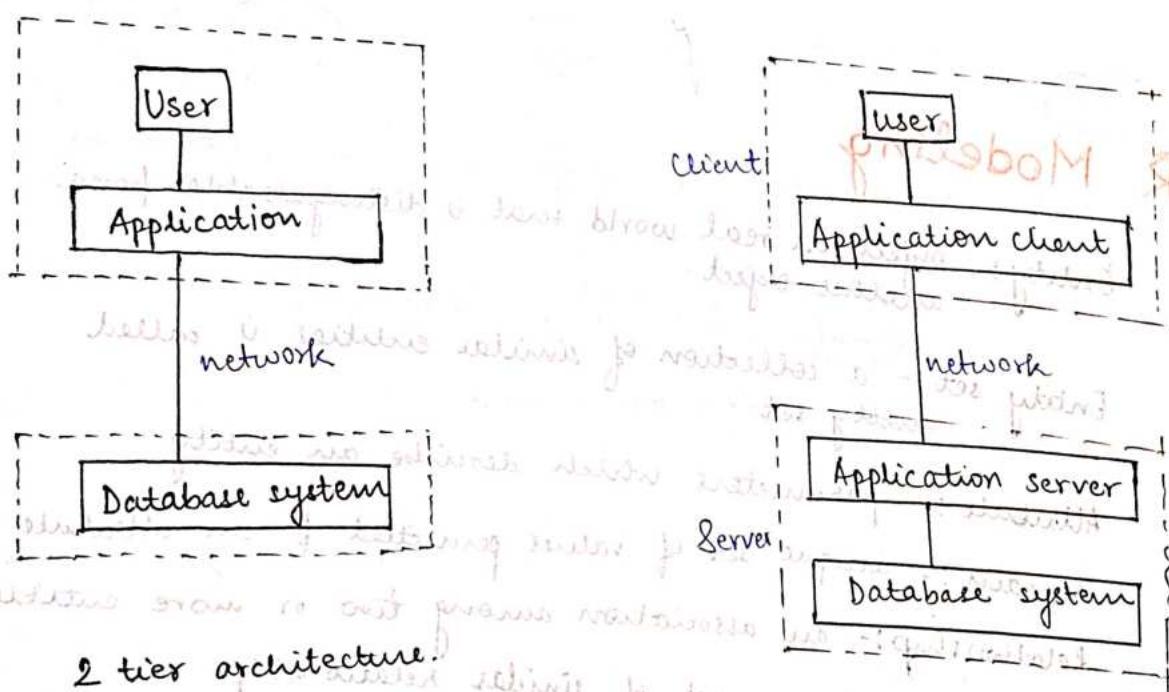
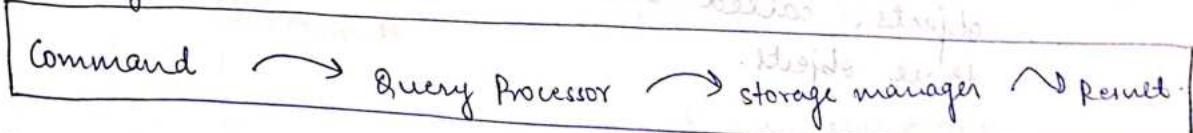
3. View level -

defines which part of database can be viewed by a person.

manager can view all the tables
employee can view only his record.



Database System Architecture.



2 tier architecture.

3 tier architecture

Data Models

A collection of conceptual tools for describing data, data relationships, data semantics, and consistency constraints.

2 models

- Entity-Relationship model (ER model)
- Relational model

1. Entity Relationship model (ER model)

ER data model consists of a collection of basic objects, called entities, and of relationships among these objects.

2. Relational model

uses a collection of tables to represent both data & relationships.

E-R Modeling

Entity :- object in real world that is distinguishable from another object.

Entity set :- a collection of similar entities is called entity set.

Attribute :- parameters which describe an entity

Domain :- unique set of values permitted for an attribute

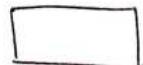
Relationship :- an association among two or more entities

Relationship-set - a set of similar relationships

Key - an attribute or a set of attributes whose values uniquely identify an entity in a set.

Descriptive attribute - attribute of a relationship.

Entity set →



Relationship →



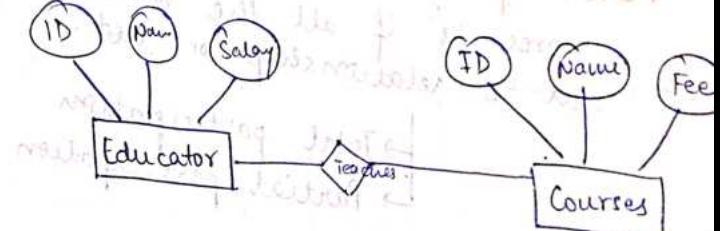
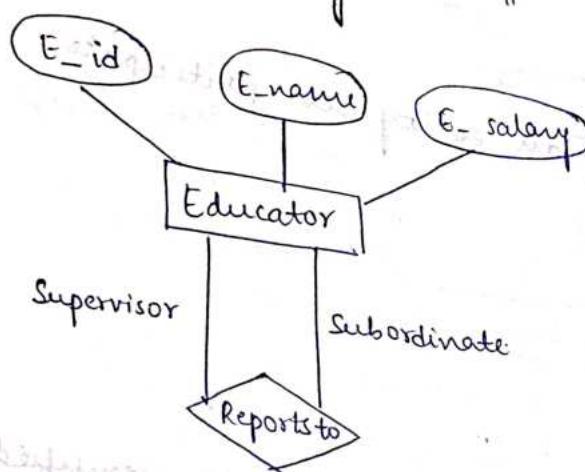
attribute →



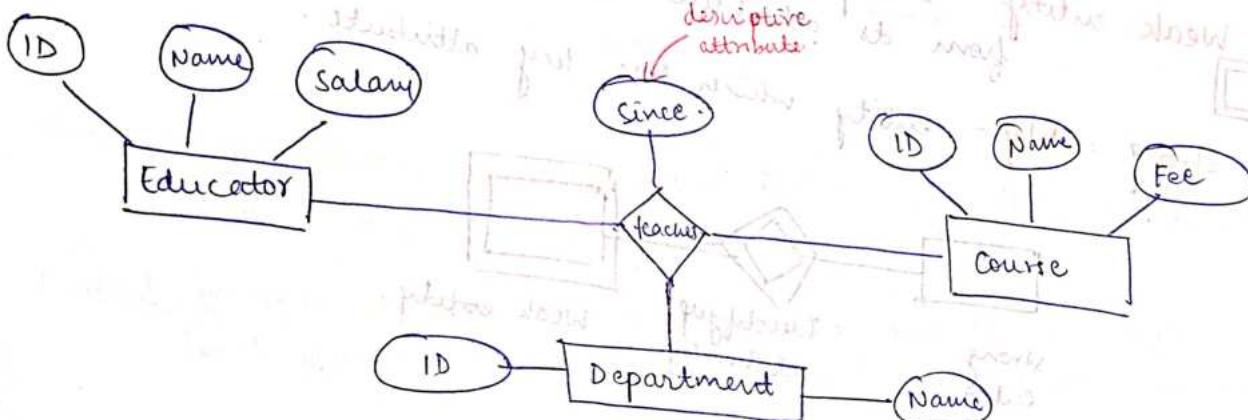
Individual assignment

Types of Relationships -

1. Unary - Relationship b/w entities of same entity set.
2. Binary - " " 2 entity sets.
3. Ternary - " " 3 entity sets.



binary relationship.



ternary relationship.

Mapping Cardinality

The maximum number of relationships in which an entity can participate is also called mapping cardinality

- One to one relationship
- One to many relationship
- Many to one relationship
- Many to many relationship

→ ~~equivalent to the right~~
→ ~~particular~~ 1
→ ~~particular~~ n
→ ~~particular~~ 2

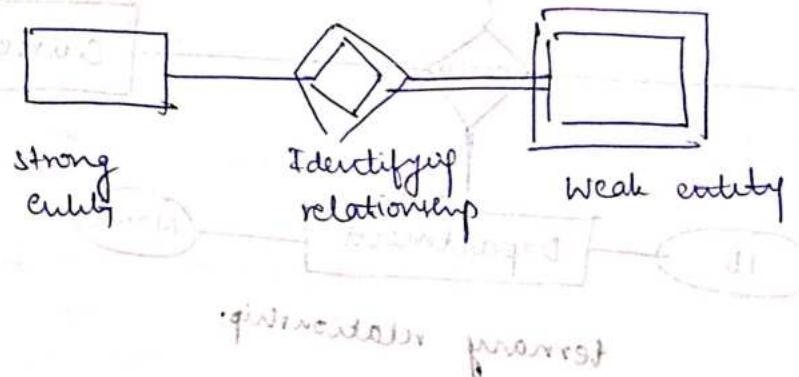
Participation constraints

represents if all the members of an entity-set participate in a relationship or not.

- Total participation
- Partial participation

Weak entity - entity which cannot be uniquely identified from its attributes alone.

Strong entity - entity which has key attribute.

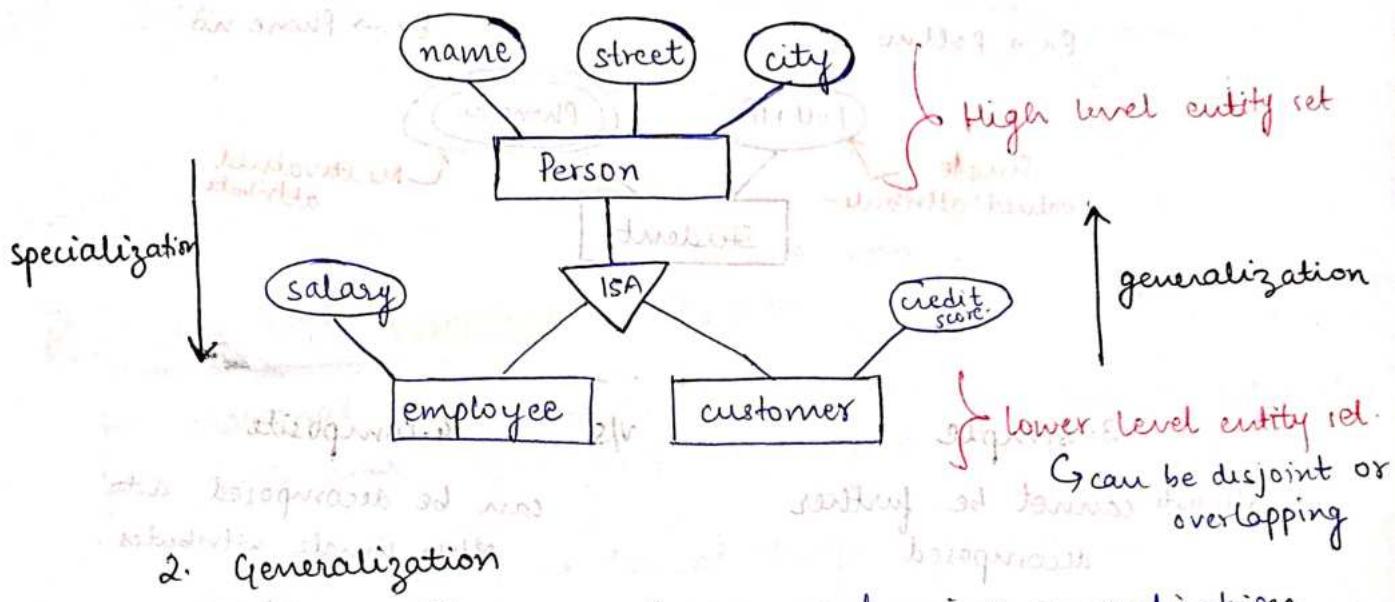


Extended ER Features

Extended ER PPT

1. Specialization

The process of designating subgroupings within an entity set is called specialization.



2. Generalization

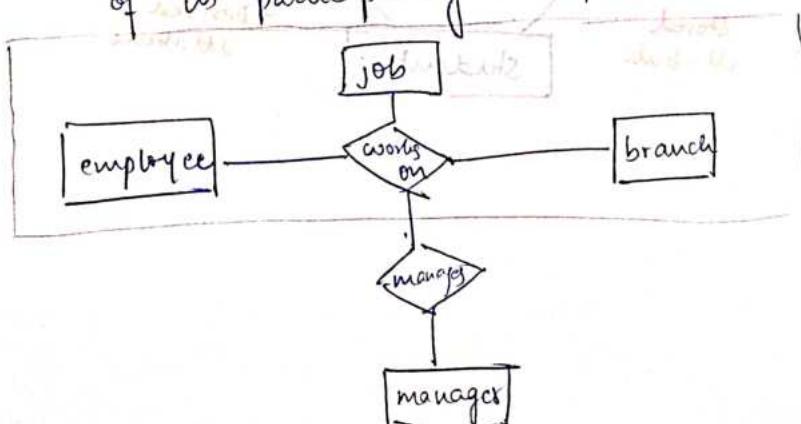
Commonality can be expressed using generalization. Contained relationship b/w higher level entity set & lower level entity set.

Total generalization / specialization - Each higher level entity set must belong to lower level entity set.

Partial generalization / specialization - Some members in higher level entity set may not belong to lower level entity set.

3. Aggregation

Relationship that contains another relationship as one of its participating entity



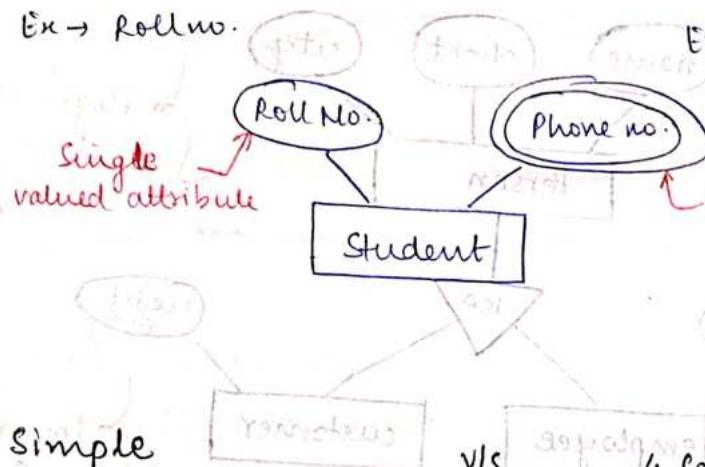
Types of attributes

22-UT-037 97 Bishash

1. Single valued

Attributes can have only one value.

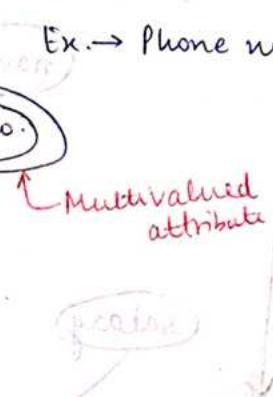
Ex. → Roll no.



2. Multi valued

Attributes can have more than one value.

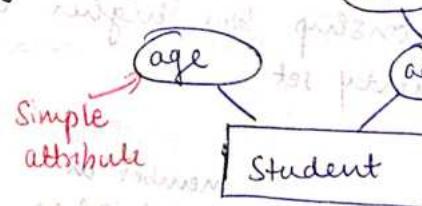
Ex. → Phone no.



3. Simple

Attribute cannot be further decomposed

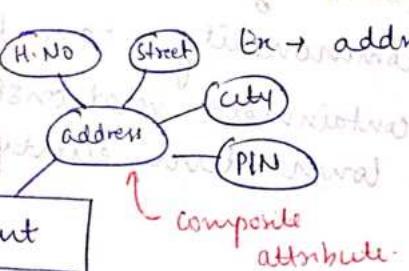
Ex. → age



4. Composite

can be decomposed into other simple attributes.

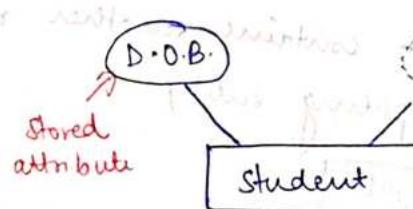
Ex. → address



5. Stored

Attributes whose value is stored in memory.

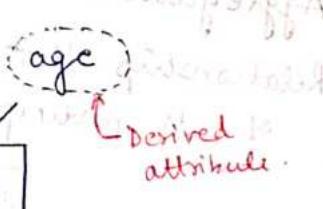
Ex. → D.O.B.



6. Derived

Attributes which can be derived from a given attribute.

Ex. → age



7. Prime attribute which is part of some key

v/s 8. Non prime attribute attribute which is not part of some key.

9. Complex attributes — Multivalued + Composite

Ex → address, phone no.

restatements of fact statements may means

Relational model

The relational model uses a collection of tables to represent both data and relationships among data

Database schema → logical design of database.

Relation_name (attribute₁, attribute₂..... attribute_n)

UNIQUE

NOT NULL ↗

Domain constraint → specifies the important condition that we want each relation to satisfy.

Degree or Arity → number of attributes in a relation

Cardinality → number of tuples in a relation

NULL → representation of no value

Key → attribute / group of attributes that can uniquely identify a tuple in a relation

Types of keys -

1. Super key - All possible keys of a relation
superkey of all keys.

2. Candidate key - Minimal superkey i.e. called candidate key.
→ A key whose proper subset is not a key.

3. Primary key -
chosen from candidate key for implementation

4. Alternate key -

any one of remaining all candidate keys other than primary key

5. Foreign key -

key referring to some attribute in other table
(all values of foreign key can be NULL)

Referential Integrity

The value in the foreign key should be present in the tab column of other table it is referring to

The foreign key should always refer to primary key

Functional Dependencies

Belief $A \rightarrow B$ (B is functionally dependent on A)
 if each value in A is associated with exactly one value
 in B in the relation.

- ① A functional dependency $\alpha \rightarrow \beta$ always holds if α is the key to the relation R .
 $\alpha, \beta \rightarrow$ attributes / set of all attributes in R .

Closure of an attribute -
 set of attributes that can be derived from the given attribute.

Ex $\rightarrow R(A, B, C, D)$

FDS $\rightarrow \{A \rightarrow B, A \rightarrow D, B \rightarrow C\}$

$$A^+ = \{A, B, C, D\}$$

$$B^+ = \{B, C\}$$

$$C^+ = \{C\}$$

$$D^+ = \{D\}$$

Ex $\rightarrow R(A, B, C, D, E)$

FDS $\rightarrow \{A \rightarrow B, B \rightarrow C, C \rightarrow D, D \rightarrow E, E\}$

$$A^+ = \{A, B, C, D, E\}$$

$$B^+ = \{A, B, C, D, E\}$$

$$C^+ = \{A, B, C, D, E\}$$

$$D^+ = \{A, B, C, D, E\}$$

$$E^+ = \{E, A, B, C, D\}$$

→ Trivial functional dependencies.

These functional dependencies are satisfied by all relations

Ex $\rightarrow A \rightarrow A$
 $\alpha \rightarrow \beta$ if $\beta \subseteq \alpha$

→ Reflexivity rule

$\alpha \rightarrow \beta$ holds if $\beta \subseteq \alpha$

→ Augmentation rule

if $\alpha \rightarrow \beta$ holds, then $\gamma\alpha \rightarrow \gamma\beta$ also holds

→ Transitivity rule

if $\alpha \rightarrow \beta$ and $\beta \rightarrow \gamma$ holds, then, $\alpha \rightarrow \gamma$ also holds.

Armstrong's axioms

Additional Rules

→ Union Rule -

If $\alpha \rightarrow \beta$ and $\alpha \rightarrow \gamma$ hold, then,
 $\alpha \rightarrow \beta\gamma$ also holds.

→ Decomposition Rule -

If $\alpha \rightarrow \beta\gamma$ holds, then,
 $\alpha \rightarrow \beta$ and $\alpha \rightarrow \gamma$ also hold.

→ Pseudo transitivity rule.

If $\alpha \rightarrow \beta$ and $\gamma\beta \rightarrow \delta$, then,
 $\alpha\gamma \rightarrow \delta$ also holds.

(A,B,C,D) ⊨ $\alpha\gamma \rightarrow \delta$

$\{A \rightarrow C, B \rightarrow A\} \models \alpha\gamma \rightarrow \delta$

Closure of a set of functional dependencies.

Finding all possible FDs from given set of FDs

Ex: R(A, B, C, D)

FDs = {A → B, B → C, AB → D}

$A^+ = A, A^+ \models F^+$

transitive rule

pseudo transitive rule

$A^+ \models F \models A^+ \models F$

$A^+ \models F \models A^+ \models F$

$A^+ \models F \models A^+ \models F$

Algorithm
for
closure

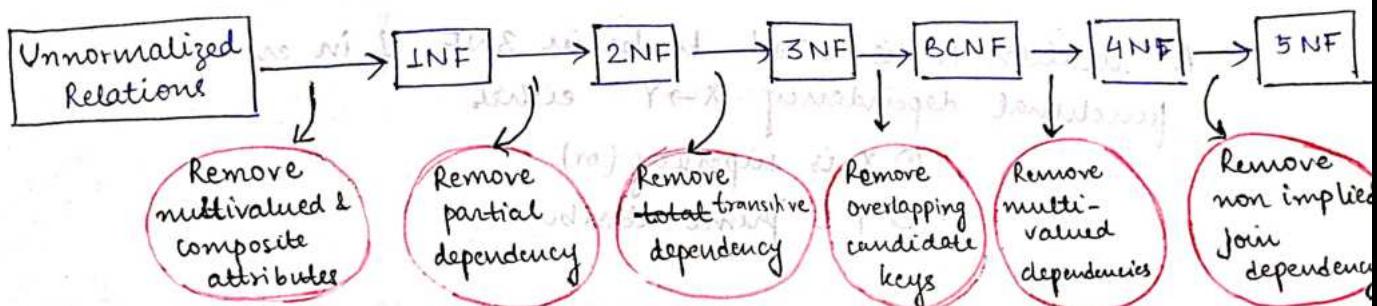
NORMALIZATION

Normalization is the process of minimizing redundancy from a relation or a set of relations.

→ can cause insert, delete, update anomalies.

- ① Process of grouping attributes into well structured relations, that contain minimum redundancy.

Process of normalization



First Normal Form (1NF)

A relation R is said to be in 1NF if there is no any multivalued & composite attribute.

Second Normal Form (2NF)

→ Relation is in 1NF

→ no ^{non} prime attribute in R should be partially dependent on R prime attribute.

Partial dependency :-

$x \rightarrow y$ is partial dependency if some attribute $A \in x$ can be removed and the dependency still holds.

→ If candidate key contains more than one attribute, then, a part of candidate key should not derive anything else.

* No partial dependency if candidate keys contains only one prime attribute.

Third normal form (3NF)

- Relation R should be in 2NF
- No any non prime attribute in R should be transitively dependent on key of R.

A relation R is said to be in 3NF if in every functional dependency $X \rightarrow Y$ either

① X is superkey (or)

② Y is prime attribute

Boyce Codd Normal Form (BCNF)

→ Relation R should be in 3NF

→ For every functional dependency $\alpha \rightarrow \beta$, α should be superkey.

* Functional dependency may not be preserved

* Any relation with 2 attributes is always into BCNF

* Normal form considered adequate for normal database design

Ans → 3NF.

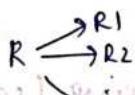
→ bcoz dependency preserving
lossless-guarantees.

Lossy Decomposition

The decomposition of relation R into $R_1, R_2, R_3, \dots, R_n$ is lossy when the join of R_1, R_2, \dots, R_n does not produce same relation as in R .

< No. of tuples are increased >

If no common attribute is present then the decomposition is lossless.



$R = R_1 \bowtie R_2 \bowtie R_3 \dots \bowtie R_n$ → Lossless Decomposition



→ Lossy Decomposition

* Dep Decomposition is lossless when the common attribute is a key in one of the sub-relations.

* Lossless decomposition is guaranteed upto 3NF.

In BCNF, sometimes, it is possible to satisfy only lossless decomposition property or functional dependency preservation property.

We focus on making decomposition lossless.

∴ Lossless decomposition holds in BCNF also.

Dependency Preserving Decomposition

A decomposition $D = \{R_1, R_2, R_3, \dots, R_n\}$ of R is dependency preserving with respect to a set F of functional dependencies if

$$(F_1 \cup F_2 \cup F_3 \dots \cup F_n)^+ = F^+$$

not ensured in BCNF

(B, 352) 338004

are derived from F and F^+ in $(B, 352)$ 338004

SQL: Structured Query Language

- ① Domain specific language
- ② used for managing data in RDBMS

→ SQL Data types

CHAR (size) → a fixed length string (can contain letters, numbers, special characters)

size → 0 < size < 255

Default → 1.

VARCHAR (size) → variable length string

size defines the max. length of the string

0 ≤ size < 65535

Introducing floating point numbers

Bit (size) → a bit value type

BOOL → 0 is considered false, non zero value is true

INT (size) → medium integer

BIGINT (size) → large integer

FLOAT (size, d) → floating point no.

d → no. of digits after decimal point.

size → total no. of digits

FLOAT(p) → floating point no.

p is used to determine whether to use FLOAT or DOUBLE

0 < p < 24 → float

25 < p < 53 → double

DOUBLE (size, d)

DECIMAL (size, d) → an exact fixed point no.

* SQL is not case sensitive

eniot_102

1. SELECT COMMAND

used to retrieve data from one or more tables.

SELECT <column-name> FROM <table-name>;

* returns
duplicates
if present.

(with values null) NOT NULL

SELECT * FROM <table-name> → returns whole table.

(with extra column)
SELECT DISTINCT <column> FROM <table-name>.

(with, aligned to) → only returns unique records
always used only with SELECT

SELECT DISTINCT country, postal FROM customers

distinct combination
of the 2 fields

NO <column> NOT DISTINCT <column> More * TELL

2. WHERE Clause

① used with * SELECT, UPDATE, DELETE & INSERT commands.

② used to filter specific rows from the table

SELECT * <column name> FROM <table-name>

WHERE <condition>

③ Relational operators that can be used in WHERE clause -

=, <>, <, <=, >, >=

④ Logical operators -

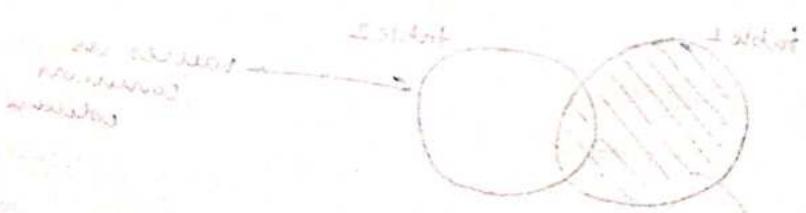
AND, OR, NOT

⑤ BETWEEN operator -

used to filter the records in specific range.

Example: <column> BETWEEN LB and UB (both inclusive)

Example: select * from student where age between 20 and 25



SQL Joins

- ① needed to retrieve records from more than one table collectively
- ② 2 tables must have one common column.

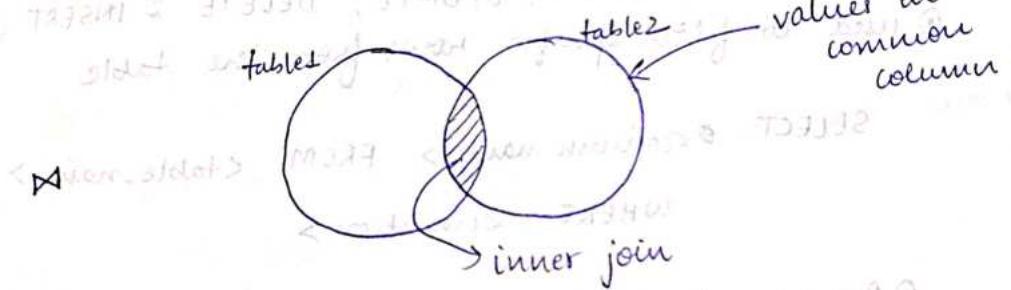
Types of joins -

- Full Join (Full outer join)
- Left Join (left outer join)
- Right Join (right outer join)
- Inner Join (Equi Join or simple join)
- Self Join

No. of attributes in result = $\text{Deg}(\text{table1}) + \text{Deg}(\text{table2})$

1. Inner join

$\text{SELECT * FROM } <\text{table1}> \text{ INNER JOIN } <\text{table2}> \text{ ON }$
 $<\text{table1}>. <\text{column}> = <\text{table2}>. <\text{column}>$



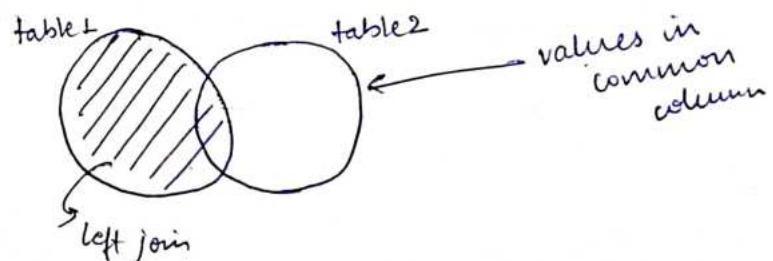
equivalent to

$\text{SELECT * FROM } <\text{table1}>, <\text{table2}> \text{ WHERE }$
 $<\text{table1}>. <\text{column}> = <\text{table2}>. <\text{column}>$

2. Left Join

$\text{SELECT * FROM } <\text{table1}> \text{ LEFT JOIN } <\text{table2}> \text{ ON }$
 $<\text{table1}>. <\text{column}> = <\text{table2}>. <\text{column}>$

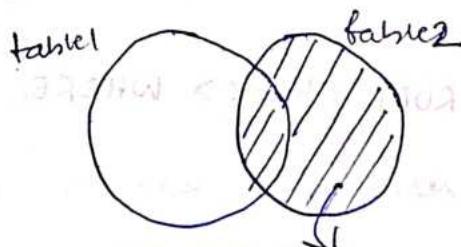
common tuples + all tuples from table1



3. Right join

SELECT * FROM <table1> RIGHT JOIN
 <table2> ON <table1>. <column> = <table2>. <column>

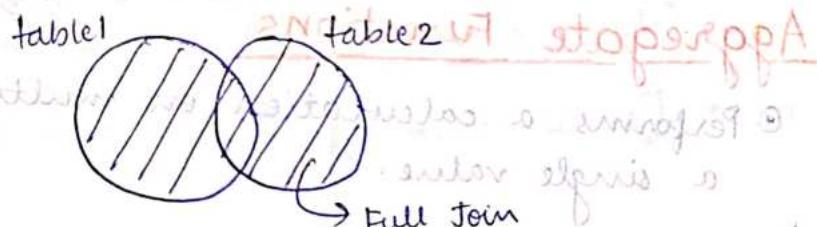
common tuples + all tuples from table 2.



4. Full join

SELECT * FROM <table1> FULL JOIN <table2>
 ON <table1>. <column> = <table2>. <column>

common tuples + tuples from table1 + tuples from table2



5. Self join

only 1 table is involved in this join

self join is not keyword

it is implemented using alias & WHERE clause

SELECT * FROM <table1> AS <t1>, <table1> AS <t2>
 WHERE <table1> <t1>. <column> = <t2>. <column>

Some figures given below show examples of self joins.

Table A

B = (A) NAME

C = (A) ADDRESS

D = A

Table B

E = (B) PVA

F = (B) DNAME

G = B

Table C

H = (C) XAM

I = (C) HNAME

J = C

Table D

K = (D) TNAME

L = (D) TNAME

M = D

Table E

N = (E) LNAME

O = (E) FNAME

P = E

Table F

Q = (F) TNAME

R = (F) LNAME

S = F

LIMIT clause

used to limit the number of fetched records from a huge database table.

$\text{SELECT * FROM <table> WHERE <condition> LIMIT <n>}$

* $\text{SELECT * FROM <table> WHERE <condition> LIMIT m, n}$

The query returns n records after leaving first m records.

leave first m records
return n records

max offset + max length + remaining offset + length

Aggregate Functions

① Performs a calculation on multiple values and returns a single value.

* ② can only be used with SELECT clause

- MIN(<column>) — returns min value in column
- MAX(<column>) — returns max. value in column
- SUM(<column>) — returns sum of values in column
- COUNT(<column>) — returns no. of records
- AVG(<column>) — returns average of values in column

all aggregate functions ignore NULL values except count(*)

A	B
1	N
2	N
N	N
3	N

$$\text{Count(*)} = 4$$

$$\text{Count(A)} = 3$$

$$\text{Count(B)} = 0$$

$$\text{SUM(A)} = 6$$

$$\text{AVG(A)} = 2$$

$$\text{MAX(A)} = 3$$

$$\text{MIN(A)} = 1$$

$$\text{SUM(B)} = \text{NULL}$$

$$\text{AVG(B)} = \text{NULL}$$

$$\text{MAX(B)} = \text{NULL}$$

$$\text{MIN(B)} = \text{NULL}$$

IN operator

- ① The IN operator allows you to specify multiple values in a WHERE clause.
- ② IN operator is shorthand for ~~not~~ multiple OR conditions.
- ③ SELECT * FROM <table> WHERE <column> IN (value1, value2...)
- ④ IN operators is also used with subqueries.

ALIAS

- ① used to give temporary name to a table, or a column of a table.
- ② only exists during duration of that query.
- ③ created using AS.
- [space is not allowed in column name of a table.
To print space, ALIAS can be used]

SELECT CustomerName AS [Customer Name] FROM Customers.

GROUP BY Clause

- ① used in collaboration with SELECT statement to arrange identical data into groups.
- ② all attributes used in GROUP BY clause must be present in the SELECT clause.
- ③ attributes not present in GROUP BY clause can appear in aggregate functions.

SELECT DNo., AVG(Salary) FROM Employees
GROUP BY DNo;

Output - (Salary of each division)

HAVING clause

similar to WHERE clause.
can only be used with GROUP BY clauses.

SELECT DNO, AVG(Salary) FROM Employees

GROUP BY DNO HAVING AVG(Salary) > 50.

Order of execution → WHERE → GROUP BY → HAVING

→ WHERE v/s HAVING

checks the condition
on individual rows

checks the condition
on a group of rows

cannot be used
with aggregate functions

can be used with
aggregate functions

executed before
GROUP BY clause

executed after GROUP
BY clause

Set operators

Set operators are used to combine results of two or more SELECT statements

no. of columns returned

from the

individual
queries must
be equal.

(SELECT * FROM Table1) <set operator> (SELECT * FROM Table2)

- UNION : (duplicate rows are eliminated)
- UNION ALL : (duplicate rows are retained)
- INTERSECT : (selects only common rows)
- MINUS / EXCEPT (set difference)

ORDER BY

- ① ORDER BY clause is used to sort the result set in ascending or descending order.

provides sorting algorithm and how works? ASC/DESC

```
SELECT * FROM <table-name> ORDER BY <column-name> ASC/DESC
                                         default
```

- ② default order of sorting → ascending order

- ③ ORDER BY clause is executed in the very end.

1) primary key
2) unique key
(Primary key = unique)

Subquery

- ① also called inner query, nested query

- ② query within a SQL query that is embedded within the WHERE clause

2 types of subquery

→ Single Row subquery

→ Multiple Row subquery.

Single Row subquery

- * Product with second highest price

SELECT MAX(price) FROM Products

WHERE price < (SELECT max(price) FROM Products)

- * Product with third highest price

SELECT MAX(price) FROM Products

WHERE price < (SELECT max(price) FROM Products)

WHERE price < (SELECT MAX(price) FROM Products);

Relational operators can be

used in single row subqueries.

Multiple Row subqueryOperators used in Multiple Row subquery1. IN operator

- Allows to specify multiple values in WHERE clause.
- Shorthand for logical OR.

`SELECT * FROM Customers WHERE country IN
(SELECT country FROM suppliers);`

* Find all customers who have placed more than 2 orders -
2 tables
Customers, Orders
`SELECT * FROM Customers WHERE customerID IN
(SELECT customerID FROM Orders
GROUP BY customerID HAVING count(customerID) > 2)`

2. ANY operator

`SELECT <column-name(s)>
FROM <table-name>
WHERE <column-name> <operator> ANY`

any standard
comparison operator

(= <> <=> < >)

3. ALL operator

`SELECT <column-name(s)>
FROM <table-name>
WHERE <column-name> <operator> ALL
(subquery)`

4. EXISTS operator

- ① used to test for the existence of any record in a subquery.
- ② returns true if subquery returns one or more records.

```
SELECT <column_name(s)>
FROM <table_name>
WHERE <column_name> EXISTS
      (subquery)
```

* 'SELECT NULL' returns 1 record with column & row NULL

null
null

RELATIONAL ALGEBRA

- Relational algebra is a procedural query language which takes instances of relations as input & yields instances of relations as output.
- provides theoretical foundation for relational databases

Basic operators—

- | | |
|-----------------------|----------------------------------|
| ① Select (σ) | ② Intersection (\cap) |
| ③ Project (π) | ④ Cartesian Product (\times) |
| ⑤ Union (\cup) | ⑥ Rename (ρ) |
| ⑦ Set Difference (-) | ⑧ Division (\div) |

1. Selection operator

- fetches tuples which satisfy a condition
- $\sigma_{\langle \text{conditions} \rangle} (\langle \text{Relation-name} \rangle)$
- similar to WHERE clause of SQL

① Selection operator is commutative

$$\sigma_{\langle \text{cond}_1 \rangle} (\sigma_{\langle \text{cond}_2 \rangle} (R)) = \sigma_{\langle \text{cond}_2 \rangle} (\sigma_{\langle \text{cond}_1 \rangle} (R))$$
$$= \sigma_{\langle \text{cond}_1 \rangle \wedge \langle \text{cond}_2 \rangle} (R)$$

2. Projection operator

- used to select only certain columns of the table

$$\pi_{\langle \text{column-name(s)} \rangle} (R)$$

* eliminates duplicates in result.

- similar to SELECT statement `SELECT DISTINCT`

RELATIONAL ALGEBRA

3. Set operations

- ① the 2 relations on which set operations are applied should necessarily have similar data types of tuples of relation
- ② they should be union compatible
 - * duplicates are not present.
 - * duplicates are present.

commutative
& associative

i) Union operator (\cup)

ii) Intersection operator (\cap)

iii) Set Difference (-)

- * $\pi_x(R_1 \cup R_2) = \pi_x(R_1) \cup \pi_x(R_2)$
- * $\pi_x(R_1 \cap R_2) \neq \pi_x(R_1) \cap \pi_x(R_2)$

4. Cartesian Product

cross product

$$|R \times S| = mn$$

$$\text{Deg}(R \times S) = \text{Deg}(R) + \text{Deg}(S)$$

$$|R| = m$$

$$|S| = n$$

→ Joins

are relations?

i) Conditional join

$R_1 \bowtie_{\text{condition}} R_2$

ii) Equi join

$$< R_1 \bowtie_{R.c = S.c} S >$$

iii) Natural join

Both relations should have atleast one common column name

$$R \bowtie S$$

iv) left outer join

$$R \bowtie L S$$

v) Right outer join

$$R \bowtie R S$$

vi) Full outer join

$$R \bowtie C S$$

5. Rename operator

① similar to aliasing in SQL.

② used to give temporary name to relation / attribute

$\rho_{\text{newname}}(\text{relation}) \rightarrow \text{Rename relation}$

$\rho_{R(A,B,C)}(R) \rightarrow \text{Rename columns to } A, B, C$

Q

$$(A \cup B) \cap (C \cup D) = (A \cap C) \cup (A \cap D) \cup (B \cap C) \cup (B \cap D)$$

6. Division operator

$$\textcircled{O} R_1 \div R_2$$

\textcircled{O} Attribute set of $R_2 \subseteq$ Attribute set of R_1 .
 Subquery minimization.

$$\text{Result set} = (\text{Attribute set of } R_1) - (\text{Attribute set of } R_2)$$

R_1	R_2	$(R_1 \div R_2) \cap R_3 = (A \cap B) \cap R_3$
$\begin{array}{ c c c } \hline A & B & C \\ \hline \end{array}$	$\begin{array}{ c c } \hline A & B \\ \hline \end{array}$	$\begin{array}{ c c } \hline A & B \\ \hline \end{array}$

not \leftarrow

$R_1 \div R_2 \rightarrow$ Those values of C which are associated with every value of A, B in R_2

<used in those queries which contain 'for all' >

Ex 9

Find value of B such that all pairs (A, B)

$\in R_1$

Find value of B such that all pairs (A, B)

$\in R_1$

Find value of B such that all pairs (A, B)

$\in R_1$

Find value of B such that all pairs (A, B)

considered as prime in listing Θ

value of B such that all pairs (A, B)

$\in R_1$

TRANSACTION

logic unit of database which includes one or more database operations

→ Commit - command used to make the changes permanent in a database. (operation successful)

→ Rollback - command used to revert back to the original state of database (operation unsuccessful)

Database state becomes same as it was on last commit

ACID Property of Transaction

1. Atomicity :- All or none. (either complete transaction or never start it)

2. Consistency :- the database state must be consistent

3. Isolation :- transactions should execute such that it does not affect nor is affected by any other concurrent transaction

4. Durability :- changes made should be reflected back into database (permanently)

Schedule

Execution ^{order} of multiple transactions on a same database concurrently is called schedule.

Why concurrency?

TRAITS

1. Improved throughput (No. of lines executed per unit time)
2. Resource utilization
3. Reduced waiting time

Problems with concurrency -

1. Recoverability problems

2. Deadlock

3. Serializability issues

Serializability problems

ACID

→ Lost update problem

→ Dirty read problem

→ Phantom tuple

→ Unrepeatable read problem

→ Incorrect summary problem.

Li least due to race between two threads

Good schedule - concurrent schedules which provide consistent result

Bad schedule - These schedules do not give consistent result (suffer from one or more of the above problems)

→ Good schedule → serializable

Serial schedules

- ① schedules which do not have interleaving
- ② always give consistent result
- ③ For n transactions, no. of serial schedules = $n!$

Non serial schedules

- ① schedules which have interleaving are called non serial schedule.

$T_1 \rightarrow n_1$ steps
 $T_2 \rightarrow n_2$ steps
 \vdots
 $T_n \rightarrow n_n$ steps

NO. of schedules possible = $\frac{(n_1 + n_2 + n_3 + \dots + n_n)!}{n_1! n_2! n_3! \dots n_n!}$

Serializable schedule

- ① A concurrent schedule which provides result as a serial schedule is called Serializable schedule.

Serializability

- conflict serializability
- View serializability

1. Conflict Serializability

→ 2 DB statements have conflict if:

- Both are in different transactions
- Both are accessing same data item
- One of them is write operation

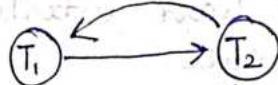
→ 2 transactions are conflict equivalent if the conflicting statements appear in the same order in both of them.

→ If one of the 2 transactions is serial schedule, the other one is called conflict serializable schedule.

To prove if a given schedule is conflict serializable,
we use precedence graph

- vertices → transactions
- edges → conflicts
- directed graph

T_1	T_2	T_3
$R(x)$		
	$W(x)$	
		$R(y)$
	$W(y)$	
$R(y)$		



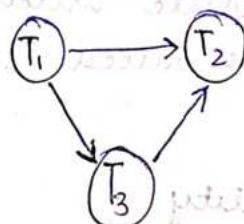
total order being maintained

as conflict is not present in precedence graph

∴ given schedule is not conflict serializable.

eg. $T_1 \rightarrow T_2 \rightarrow T_3$

T_1	T_2	T_3
$R(x)$		
	$R(y)$	
		$R(y)$
	$W(y)$	
		$W(x)$
	$R(x)$	
	$R(x)$	



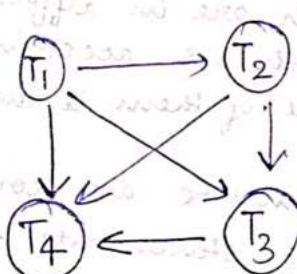
admits interleaving?

no cycle

∴ conflict serializable schedule

eg. serial schedule = T_1, T_3, T_2

T_1	T_2	T_3	T_4
$R(A)$			
$W(B)$	$R(B)$		
		$R(C)$	
$W(A)$			$R(A)$
			$W(A)$
	$R(C)$		
			$W(B)$
		$R(B)$	
			$W(C)$



conflict serializable schedule

eg. conflict serial schedule

eg. serial schedule

T_1	T_3	T_4	T_2
T_1	T_3	T_2	T_4
T_1	T_2	T_3	T_4
T_2	T_1	T_3	T_4

2. View Serializability

Because of view

3 points to remember

↳ who reads from database

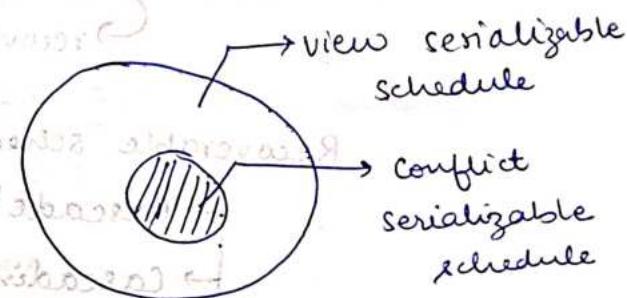
↳ who reads from other's written value

↳ who writes last

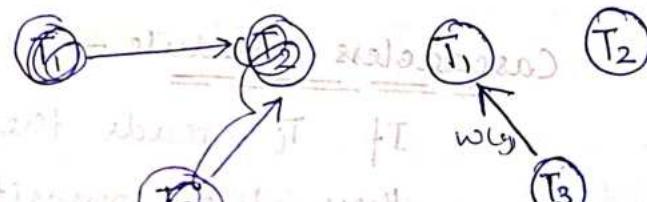
② 2 schedules s_1 , s_2 are view equivalent if both are following similar sequence/rules of view serializability.

① All conflict serializable

schedules are view serializable.



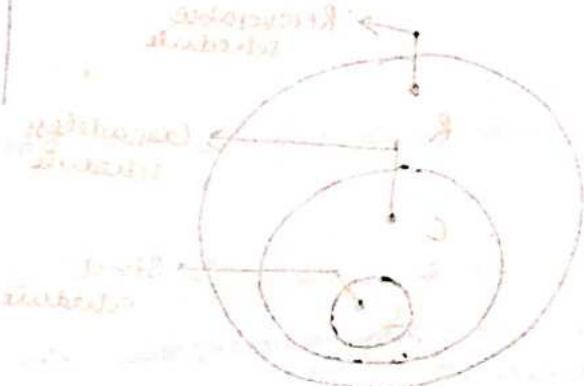
	T ₁	T ₂	T ₃
R(X)			
	R(Y)		
		w(Y)	
w(Y)			
			w(X)
w(X)			



not view serializable

T ₃	T ₁	T ₂
T ₃	T ₁	T ₂
	T ₂	T ₁

3 possible serial schedules



Recoverability

When no any committed transaction needs to roll back, it is called recoverable schedule.

If a transaction T_i reads dirty value written by T_j and then, the commit operation of T_i should be after the transaction T_j commits.

Recoverable schedule

Recoverable schedule

→ Cascadeless Recoverable schedule Rollback

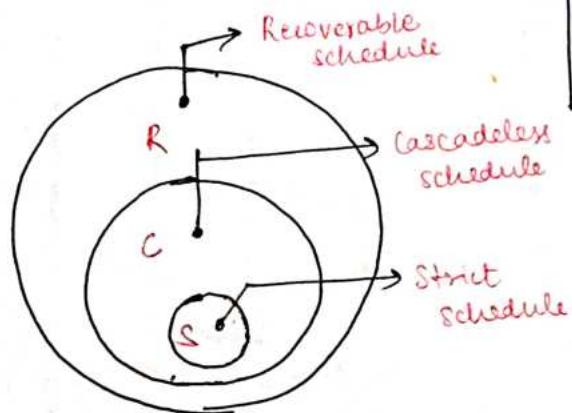
→ Cascading Recoverable schedule Rollback

($\sqsubset T$) Cascadeless schedule -

If T_i reads the data item written by T_j then, the commit operation of T_j should appear before read operation of T_i .

T_j	T_i
$R(x)$	
$x = x + 2$	
$w(x)$	
commit	

$R(x)$
 $x = x + 3$
 $w(x)$
commit



Obtaining serializability (view serializable & conflict serializable) is also not practically possible.
∴ locking protocols are used.

Locking Protocols

- Lock  Shared lock (other transactions can share)
taken for read
- Lock  Exclusive lock (no other transaction can read/write)
taken for write.

Demand	Shared	Exclusive
Shared	✓	✗
Exclusive	✗	✗

* A transaction may suffer from starvation if it demands exclusive lock while shared lock is being held by multiple transactions.

low to avoid starvation
 * If a transaction is blocked for exclusive lock because other transactions have shared lock, then multiple shared locks are not allowed.

Table maintained by DBMS	Availability	Mode	Count	Blocked
X	0 if available 1 if lock is held.	S → shared EX → Exclusive	No. of transactions holding lock (only in case of shared lock)	Transactions (FIFO order) that are waiting for the data item to be unlocked.

Simple locking mechanism can lead to unrepeatable read problem.

T ₁	T ₂
Lock-S(x) RLX Unlock(x)	LOCK-X(x) W(x) Unlock(x)

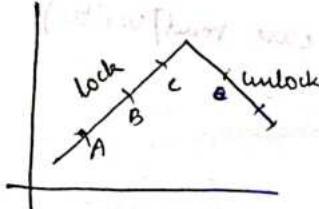
∴ To solve unrepeatable read problem, several protocols / rules are developed.

2 Phase Locking Protocols (Basic 2PL)

- Systematic locking mechanism.
- Once unlock is done, a transaction is not allowed to lock any other data item.

Each transaction has 2 phases -

- Growing phase - acquiring of locks
- Shrinking phase - releasing of locks.



T_1	T_2
LOCK-S(x)	
R(x)	
LOCK-EX(y)	
Unlock(x)	
	LOCK-EX(x)
	W(x)
W(y)	
Unlock(y)	

∴ Only 2 possible schedules are there -

T_1	T_2
R(x)	
	W(x)
	W(y)
LOCK-S(y)	
R(y)	
Unlock(x)	
Unlock(y)	

T_1	T_2
R(x)	
	W(y)
W(x)	
R(y)	

T_1	T_2
R(x)	
	W(x)
W(y)	
R(y)	

schedule not allowed by basic 2PL.

Every schedule which is allowed under basic 2PL is conflict serializable also.

Question asked - Tell whether the given schedule is allowed under 2PL or not.

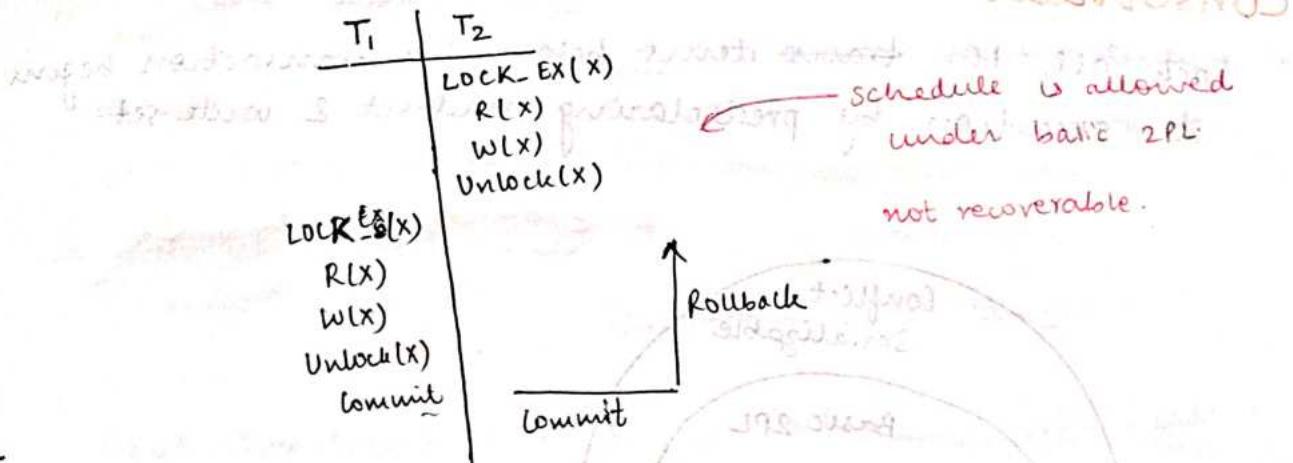
T_1	T_2	T_3	T_4
W(x)			
	W(y)		
		W(z)	
			R(x)
			R(y)
			R(z)

* 2PL suffers from starvation.

* 2PL suffers from deadlock.

T_1	T_2
LOCK-S(x)	
R(x)	
	LOCK-EX(y)
	W(y)
	LOCK-EX(x)
	W(x)

* 2 PL can allow non recoverable schedules



Strict Schedule - Every write operation should be ended before another transaction performs read/write on the same database item.

Strict 2PL Protocol

Exclusive locks should be released only after commit.

- * Strict 2PL allows only strict schedules.

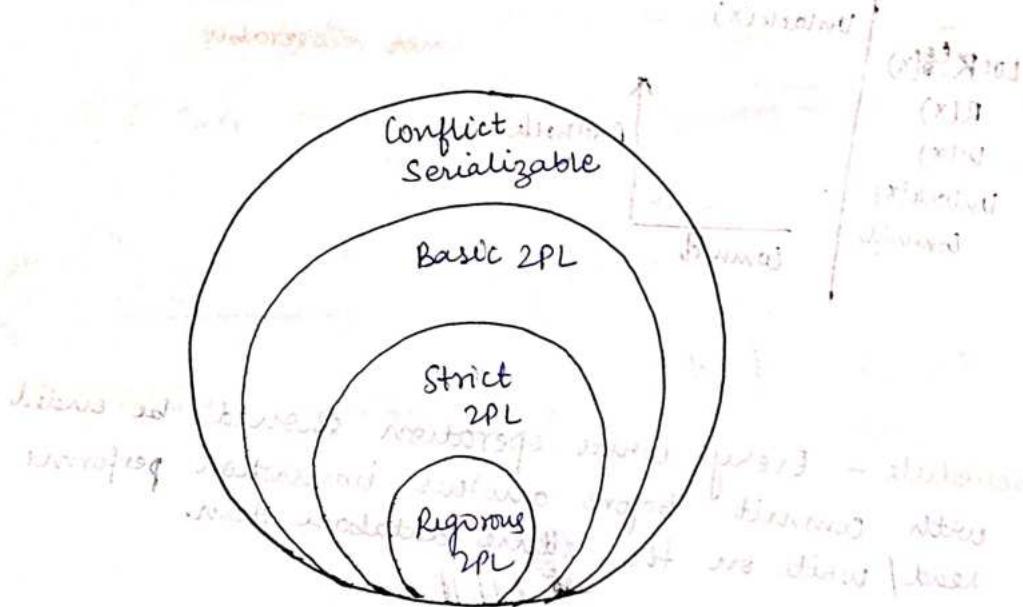
Rigorous 2PL Protocol

Every lock should be released after commit.

- ① Strict & Rigorous 2PL do not suffer from dirty read problem.
- ② Basic 2PL may allow non recoverable schedule.
- ③ Strict & Rigorous 2PL have only recoverable schedules (cascadelock)

Conservative 2PL - Protocol

Lock all the trans items before the transaction begins its execution by predeclaring read-set & write-set



Same as 2PL plus weaker as it doesn't allow missing

first two are serializing
all weaker forms

allowing 2PL warping
allowing weaker as it doesn't allow missing

last point may suffice for ab 2PL warping 2PLs @

strictly stronger than warping form 2PLs @

strictly weaker than 2PL warping 2PLs @

(strictly) weaker

Timestamp Based Protocols

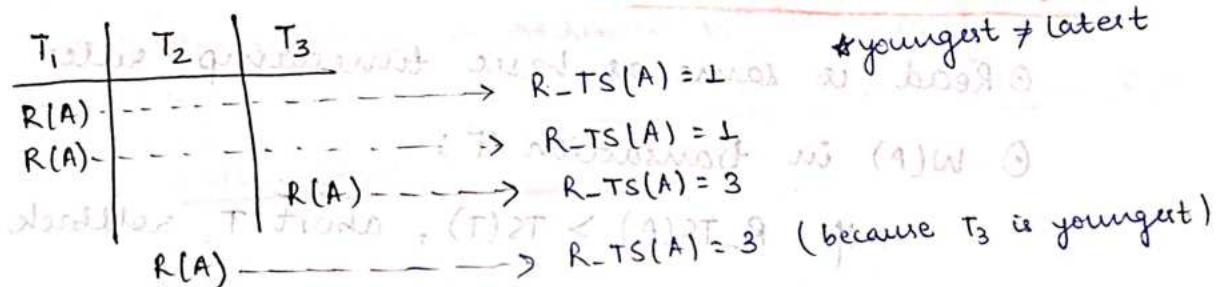
Timestamp:- unique identifier created by DBMS to identify the relative starting time of a transaction

$TS(T_1) < TS(T_2) \rightarrow T_1 \text{ starts before } T_2$

$(T_1 \text{ is older than } T_2)$

Read Timestamp(A) \rightarrow Youngest transaction who reads A

Write Timestamp(A) \rightarrow Youngest transaction who writes A



- ① Basic Timestamp algorithm is used to control concurrency by allowing transactions to run in their relative ordering of timestamp.



When a transaction T issues $w(x)$ operation, check the following conditions -

- ① If $R_TS(x) > TS(T)$ or $W_TS(x) > TS(T)$ abort the transaction and rollback T .

- ② else, execute $w(x)$ and update $W_TS(x)$ to $TS(T)$.

When transaction T issues $R(X)$ operation, check the following conditions -

- ① If $W_{TS}(X) > TS(T)$ then abort and reject the operation.
- ② else execute $\theta R(X)$ and set $R_{TS}(X)$ to $\max(TS(T), R_{TS}(X))$

Thomas Write Rule

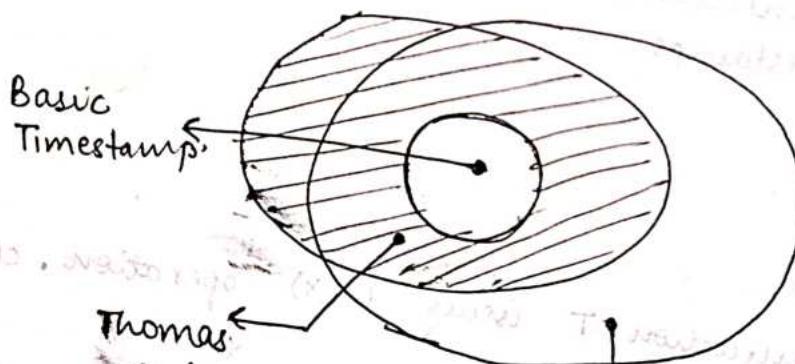
① Read is same as basic timestamp rules.

② $W(A)$ in transaction T:

→ If $R_{TS}(A) > TS(T)$, abort T, rollback

→ else if $W_{TS}(A) > TS(T)$, skip the write operation

→ else perform $W(A)$ and update $W_{TS}(A) = TS(T)$

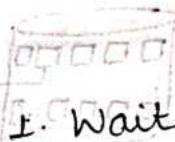


- ① Some schedules which follow Thomas write rule are not conflict serializable

Deadlock Prevention

1. Wait-DT

Let T_i = older transaction T_j = newer younger transaction
 $TS(T_i) < TS(T_j)$



1. Wait Die -

→ If T_i tries to acquire lock currently held by T_j , then, T_i waits till the lock is available.

→ If T_j tries to acquire lock currently held by T_i , then, T_j is aborted and restarted with same time stamp.

2. Wait Wound -

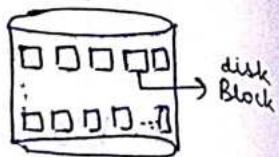
→ If T_i tries to acquire lock currently held by T_j , then, T_j is pre-empted by aborting it.

→ If T_j tries to acquire lock currently held by T_i , then, T_j waits till the lock is available.

Time	Lock	Owner	Lock Type
1	L1	T_1	R
2	L1	T_2	R
3	L2	T_1	R

Disk Block

- ① logical unit of storage created by file system.
- ② A combination of 1-2 sectors.
- ③ size = 512B to 2KB



bits of information stored at least $512 \times 8 = 4096$ bits

Record storage

Suppose, 1 disk block = 64 bytes = 2^6 bytes

bits of information stored at least $64 \times 8 = 512$ bits

Illustrating

Name	RollNo.	...

→ 1 record = 4 bytes
 \therefore No. of records in 1 disk block = $\frac{2^6}{2^2} = 2^4 = 16$

If the table contains 512 records,

No. of disk blocks needed = $\frac{512}{16} = 32$ blocks.

bits of information stored at least $32 \times 64 = 2048$ bits

Fixed v/s Variable Length Records

bits of information stored at least $512 \times 8 = 4096$ bits

student table

Roll No.	Name	Photo (max. 2 MB)	Signature (max. 200 KB)

max. size of 1 record ≈ 2.2 MB

If a student uploads pic of size 1 MB, sign 200 KB,

size of record = 1.2 MB.

In fixed length record, size of record is fixed = 2.2 MB
(extra space is wasted)

In variable length record, size of record is not fixed
(space ~~wasted~~ utilized)

Spanned Mapping

record of a file is stored inside the block even if it can only be stored partially, hence, record is spanned over multiple blocks.



Adv :- no wastage of memory
(internal fragmentation)



Disadv :- memory access time
is larger.

Block size = 32B

Record size = 5B

6 records in block 1

2B of 7th record in b1

3B of 7th record in b2

Unspanned Mapping

a record of file is stored inside a block only if it can be completely stored inside it.



Adv :- access time of records is less



Disadv :- wastage of memory
(internal fragmentation)

<space wasted per block>

Block size = 32B

Record size = 5B

6 records in B1

7th records in B2

2B in B1 wasted

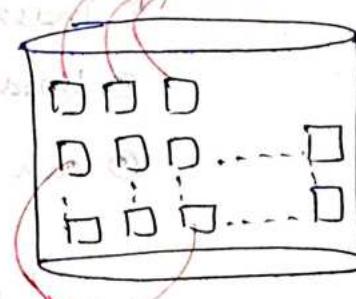
Indexing

index file is created to store information of which records are present in which disk block.

Index file/table is also stored in disk blocks.

→ Dense index:

Index record is for each database record.



→ Sparse Index:

Index record will for a few few records only.

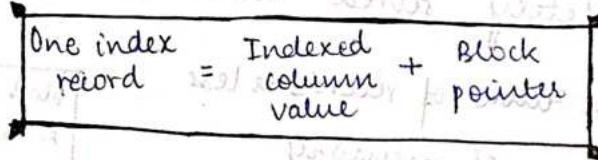
DB records.

Indexing Techniques

- Primary indexing
- clustering indexing
- Secondary key indexing
- Secondary non key indexing

1. Primary indexing

- ① Indexing is done on key primary key or super key.
- ② Data must be ordered on index.
- ③ It's always sparse index



2. Clustering indexing

- ① Indexing is done on non key field. (duplicate values may be present)
- ② Data must be ordered on index field

→ Indexing is done for unique value of non key

③ sparse indexing

3. Secondary key indexing

- ① Indexing done on primary key or any super key
- ② Data must not be ordered on index field
- ③ can be dense or sparse

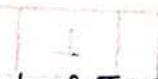
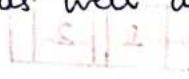
<non key> → similar to key indexing
except that indexing is done on non key field.

B Tree

- aligns E

- ① Tree based indexing technique.
- ② Dynamic indexing technique.
- ③ Self balancing tree search tree

→ grows horizontally as well as vertically.



Order of B tree - max children of a node of B Tree
(order is denoted by 'p')

An order p B-Tree -

↳ every node has max p children (tree pointers)

↳ every node other than root should have atleast $\lceil \frac{p}{2} - 1 \rceil$ nodes.

↳ every node can have almost $(p-1)$ keys and p tree pointers

↳ Root can have minimum 1 node key

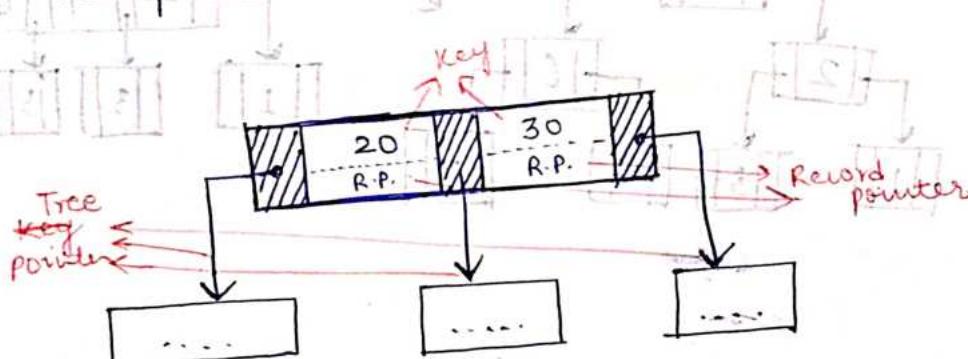
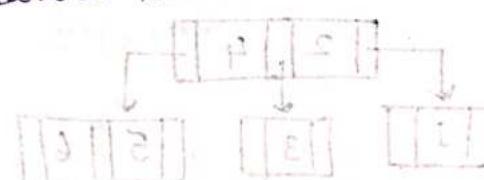
↳ all leaves are at same level.

B Tree node structure

Key - index pointer

Record pointer :- pointer to the disc block where key index is stored.

Tree pointer :- children



Example -

order 3 B Tree

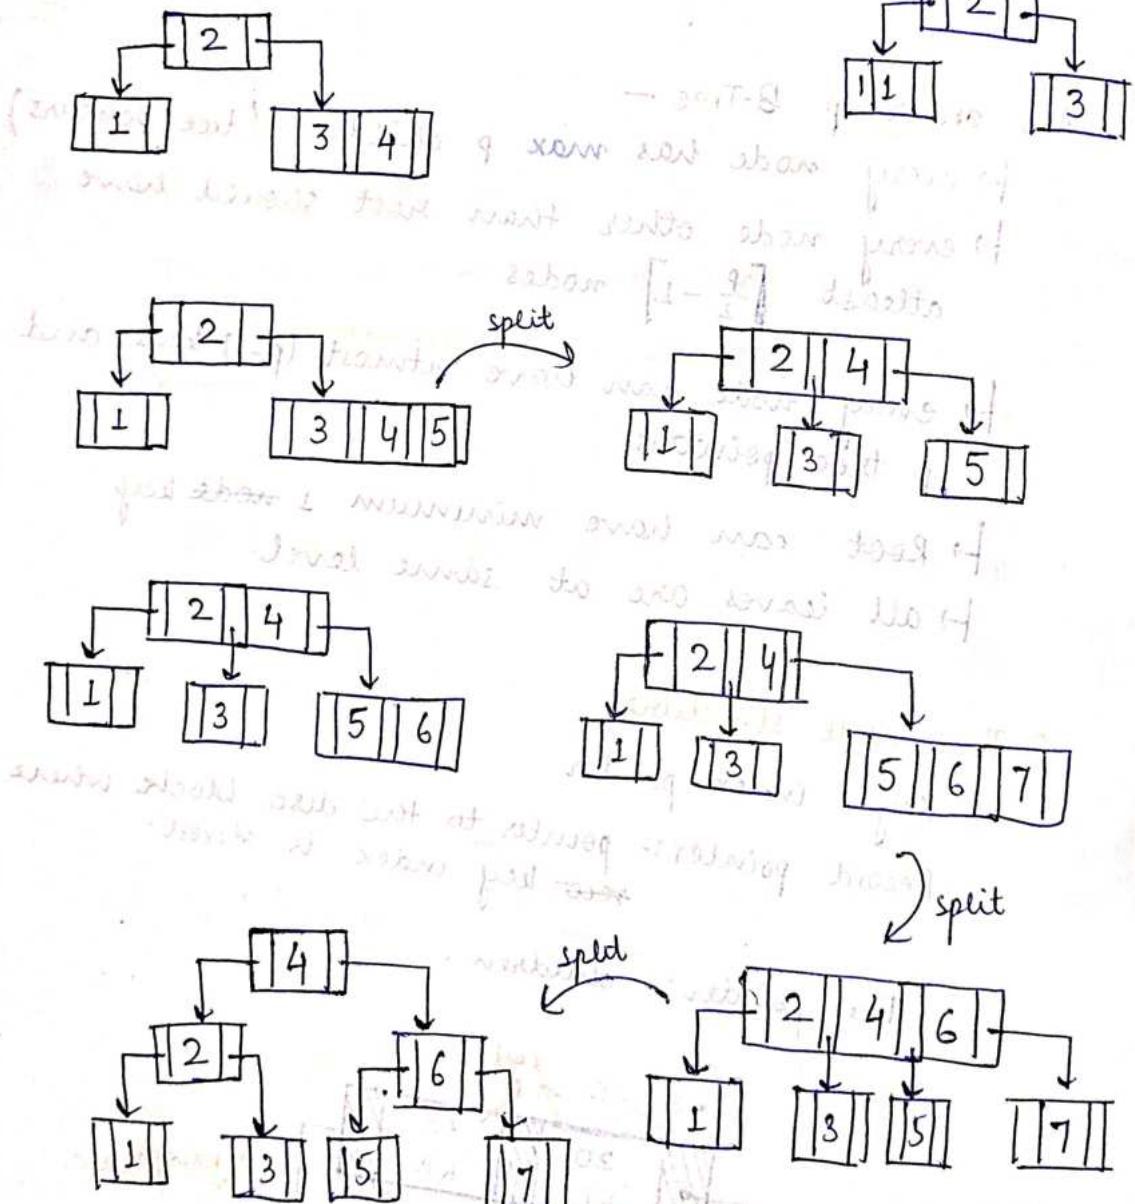
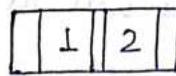
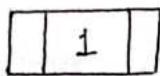
Max keys = 2

Max links (children) = 3

$$\text{Min keys (except root)} = \lceil \frac{3}{2} \rceil - 1 = 1$$

$$\text{Max keys (except root)} = \lceil \frac{3}{2} \rceil = 2$$

Insert keys 1, 2, 3, 4, 5, 6, 7



Each node of a B Tree is stored in a disk block.

If p is order of tree,

$$\text{max No. of keys} = p-1$$

$$\text{max. No. of record pointers} = p-1$$

$$\text{max. No. of tree pointers} = p.$$

$$\therefore (p-1)(\text{key size}) + (p-1)(\text{record pointer size}) + p(\text{tree pointer size}) \leq \text{Block size}$$

formula used to find

max. value of ' p '

If block size is given.

or a min block size if p is given.

Inp:-

p order B tree

Total nodes = n

$$H_{\min} = \text{min height of BTree} = \lceil \log_p(n+1) - 1 \rceil$$

$$H_{\max} = \text{max height of BTree} = \lfloor \log_{\lceil \frac{p}{2} \rceil} \frac{n+1}{2} \rfloor$$

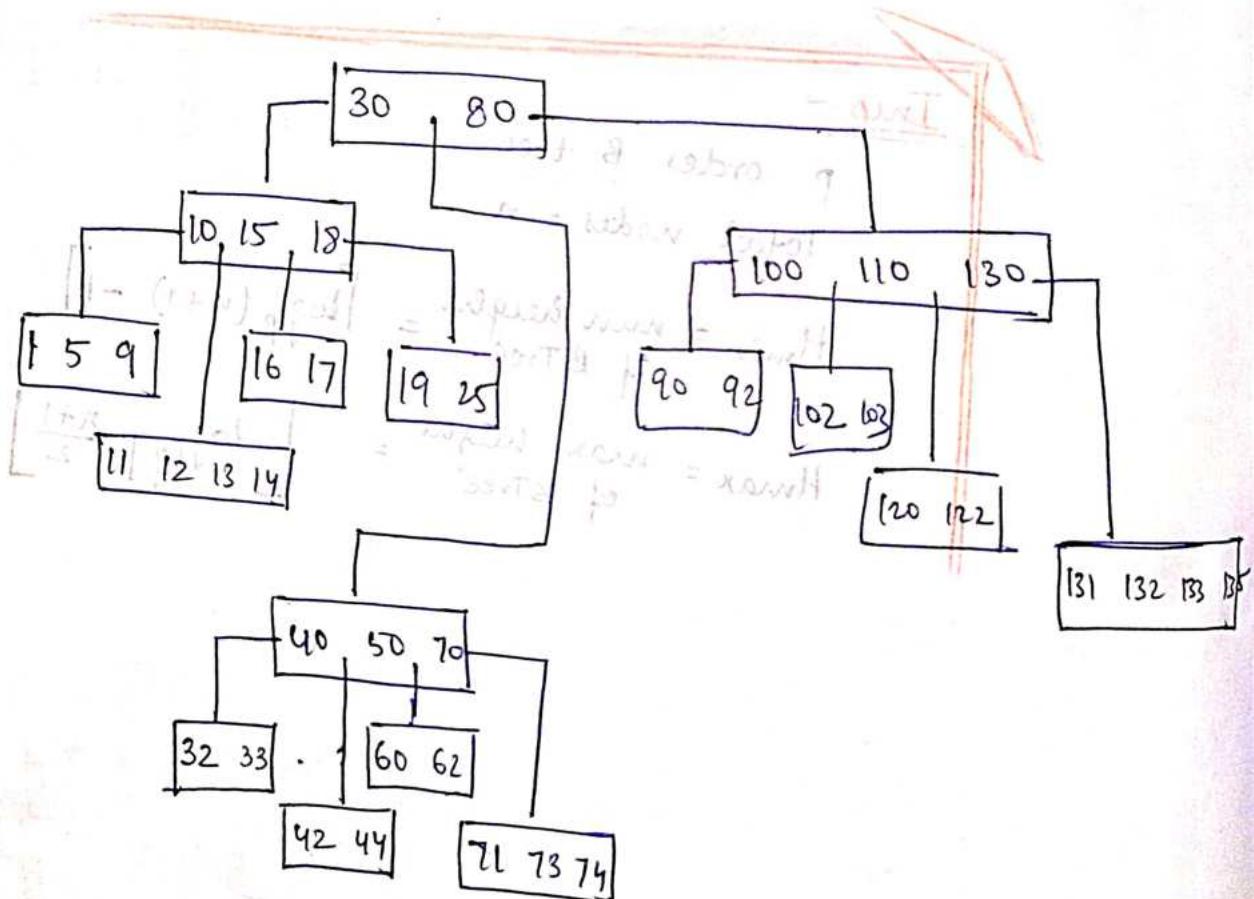


- B-Tree deletion in leaf node -
- After deletion, if no violation of new minimum keys then no changes in tree.
 - If violation of min keys, then, borrow key from sibling (rotation through parent)
 - If borrow through parent can't be possible then, merge the node with sibling & pull down the ~~set anchor node~~ key from parent.

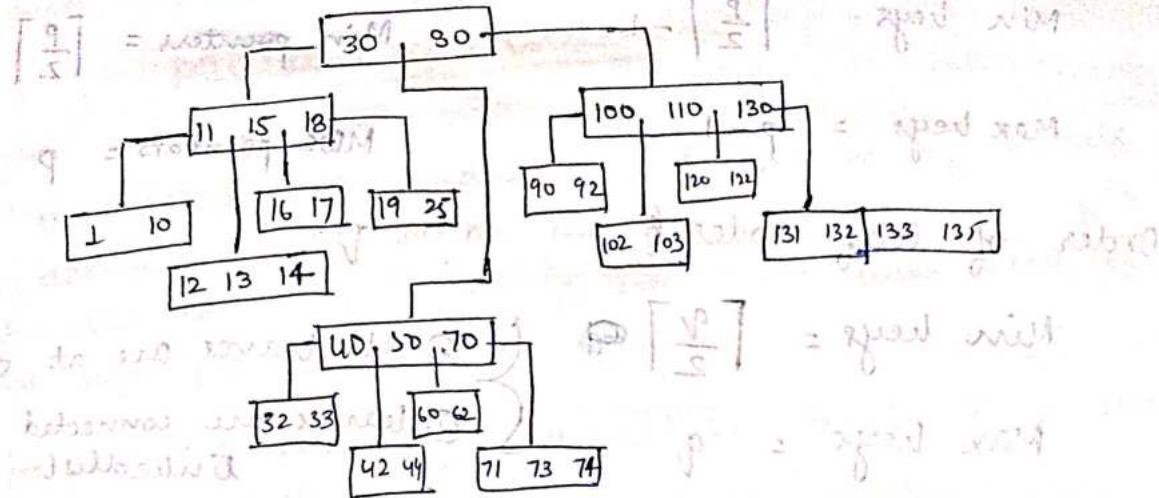
• Btree of order 5

$$\text{max keys} = 4$$

$$\text{min keys} = \left\lceil \frac{5}{2} \right\rceil - 1 = 2$$



Delete 5 → no change
Delete 9 → borrow from sibling → transfer from left to right



B tree deletion from internal node -

→ Replace the deleted value with inorder successor or inorder predecessor.

→ Now follow rules of deletion of key from leaf node.

B⁺ Tree

Internal nodes

Keys
Tree pointers

order of leaf node and internal node

leaf nodes

Keys
Record pointers
Linked list pointer

is different

○ Leaf nodes are connected to each other or linked list

○ Range queries work faster in B⁺ Trees.

Order of internal nodes (except root) \rightarrow op order P.

$$\text{Min keys} = \left\lceil \frac{P}{2} \right\rceil - 1 \quad \text{Min pointers} = \left\lceil \frac{P}{2} \right\rceil$$

$$\text{Max keys} = p-1 \quad \text{Max pointers} = P$$

Order of leaf nodes \rightarrow order q

$$\text{Min keys} = \left\lceil \frac{q}{2} \right\rceil \quad \begin{cases} \textcircled{1} \text{ all leaves are at same level} \\ \textcircled{2} \text{ leaves are connected using linkedlist.} \end{cases}$$

$$\text{Max keys} = q$$

- * Splitting at internal node \rightarrow same as that of B Tree
- splitting at leaf node - middle key is duplicated & is part of right subtree.

above fool work just for waitlisted for more waiting areas

① One node of B+ tree must be present in one disk block.

$$\text{order of B+ tree} = p$$

Key = 8 bytes

Block pointer = 32 bytes

Record pointer = 48 bytes

disk block size = 2048 bytes

For internal node,

$$\text{max } (p-1) \text{ keys}$$

~~BPT + 8~~
p block pointers

$$\therefore (p-1)8 + p(32) \leq 2048$$

$$\Rightarrow 8p - 8 + 32p \leq 2048$$

$$\Rightarrow 40p \leq 2048 + 8 \leq 2056$$

$$\Rightarrow p \leq \frac{2046}{40} = 51.125$$

For external node,

$$\text{max } (q-1) \text{ keys}$$

~~+(q-1) record pointers~~

$$\therefore p = 51 \text{ (max)}$$

~~from max order of B+ tree = 51~~

~~internal node~~

~~max order of leaf node~~

$$\therefore (q-1)(8) + (q-1)(48) \leq 2048$$

$$\Rightarrow (q-1)(56) \leq 2048$$

$$\Rightarrow q-1 \leq \frac{2048}{56}$$

$$\Rightarrow q-1 \leq 36.57 \Rightarrow q \leq 37.57$$

$$\therefore q = 37 \text{ (max)}$$

max order of leaf node
of B+ tree = 37.

Multip

Multi-level indexing

- In secondary key indexing, if the no. of records in the database table is huge, then, the index file size will also be very huge & searching will need to access multiple disk blocks.
 - entire index file cannot be stored in a single disk block.
 - ∴ Do indexing on the index file.
- < Primary indexing >

No. of blocks accessed needed to search a record is minimized.

